SF Salaries Exercise

Welcome to a quick exercise for you to practice your pandas skills! We will be using the SF Salaries Dataset from Kaggle! Just follow along and complete the tasks outlined in bold below. The tasks will get harder and harder as you go along.

** Import pandas as pd.**

```
import pandas as pd
```

** Read Salaries.csv as a dataframe called sal.**

```
df=pd.read_csv("salaries1.csv")
```

** Check the head of the DataFrame. **

```
df.head()
```

	Id	EmployeeName	JobTitle	BasePay	OvertimePay	OtherPay	Benefits	Tota
0	1	NATHANIEL FORD	GENERAL MANAGER-METROPOLITAN TRANSIT AUTHORITY	167411.18	0.00	400184.25	NaN	56759
1	2	GARY JIMENEZ	CAPTAIN III (POLICE DEPARTMENT)	155966.02	245131.88	137811.38	NaN	53890
2	3	ALBERT PARDINI	CAPTAIN III (POLICE DEPARTMENT)	212739.13	106088.18	16452.60	NaN	33527
			WIRE ROPE					

** Use the .info() method to find out how many entries there are.**

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 148654 entries, 0 to 148653
Data columns (total 13 columns):
 #   Column            Non-Null Count   Dtype  
---  ------            --------------   -----  
 0   Id                148654 non-null  int64  
 1   EmployeeName      148654 non-null  object 
 2   JobTitle          148654 non-null  object 
 3   BasePay           148045 non-null  float64
 4   OvertimePay       148650 non-null  float64
 5   OtherPay          148650 non-null  float64
 6   Benefits          112491 non-null  float64
 7   TotalPay          148654 non-null  float64
 8   TotalPayBenefits  148654 non-null  float64
 9   Year              148654 non-null  int64  
 10  Notes             0 non-null       float64
 11  Agency            148654 non-null  object 
 12  Status            0 non-null       float64
dtypes: float64(8), int64(2), object(3)
memory usage: 14.7+ MB
```

What is the average BasePay ?

```
df["BasePay"].mean()
```

```
66325.44884050643
```

** What is the highest amount of OvertimePay in the dataset ? **

```
df["OvertimePay"].max()
```

```
245131.88
```

** What is the job title of JOSEPH DRISCOLL ? Note: Use all caps, otherwise you may get an answer that doesn't match up (there is also a lowercase Joseph Driscoll). **

```
df[df["EmployeeName"]=="JOSEPH DRISCOLL"]["JobTitle"]
```

```
24    CAPTAIN, FIRE SUPPRESSION
Name: JobTitle, dtype: object
```

** How much does JOSEPH DRISCOLL make (including benefits)? **

```
df[df["EmployeeName"]=="JOSEPH DRISCOLL"]["TotalPayBenefits"]
```

```
24    270324.91
Name: TotalPayBenefits, dtype: float64
```

** What is the name of highest paid person (including benefits)?**

```
df[df["TotalPayBenefits"]==df["TotalPayBenefits"].max()]
```

	Id	EmployeeName	JobTitle	BasePay	OvertimePay	OtherPay	Benefits	TotalI
		NATHANIEL	GENERAL MANAGER-					

** What is the name of lowest paid person (including benefits)? Do you notice something strange about how much he or she is paid?**

```
df[df["TotalPayBenefits"]==df["TotalPayBenefits"].min()]
```

	Id	EmployeeName	JobTitle	BasePay	OvertimePay	OtherPay	Benefits	Total
			Counselor,					

** What was the average (mean) BasePay of all employees per year? (2011-2014) ? **

```
df.groupby("Year").mean()["BasePay"]
```

```
Year
2011    63595.956517
2012    65436.406857
2013    69630.030216
2014    66564.421924
Name: BasePay, dtype: float64
```

** How many unique job titles are there? **

```
df["JobTitle"].nunique()
```

```
2159
```

** What are the top 5 most common jobs? **

```
jobs=df.groupby("JobTitle").count()
top=jobs.sort_values(by="Id", ascending=False)[:5]
top["Id"]
```

```
JobTitle
Transit Operator                7036
Special Nurse                   4389
Registered Nurse                3736
Public Svc Aide-Public Works    2518
Police Officer 3                2421
Name: Id, dtype: int64
```

** How many Job Titles were represented by only one person in 2013? (e.g. Job Titles with only one occurence in 2013?) **

```
year=df[df["Year"]==2013]
group=year.groupby("JobTitle").count()
count=group[group["Id"]==1]
count.count()["Id"]
```

```
202
```

** How many people have the word Chief in their job title? (This is pretty tricky) **

```
def fun(job_title):
    if "chief" in job_title.lower().split():
        return True
    else:
        return False
df=pd.read_csv("salaries1.csv")
sum(df["JobTitle"].apply(lambda x: fun(x)))
```

```
477
```

** Bonus: Is there a correlation between length of the Job Title string and Salary? **

```
df["title_len"]=df["JobTitle"].apply(len)
df[["title_len","TotalPayBenefits"]].corr()
```

	title_len	TotalPayBenefits
title_len	1.000000	-0.036878
TotalPayBenefits	-0.036878	1.000000

Great Job!